


ETHICAL LESSONS LEARNED FROM COMPUTER SCIENCE



by Richard Bergmair

Introduction

In this article, we will address the question “How can computer science methods help us to better understand ethics?”

Ethics, and especially normative ethics, are concerned with finding “criteria for what is morally right and wrong. Ethics includes the formulation of moral rules” [1].

Logic programming is concerned with identifying base facts that are logically true or false. It includes the formulation of inference rules.

The polymorphy described above sits at the core of this discussion. To show this more clearly, a type of logic will be established, but instead of reasoning, and determining whether something is true or false, it will make moral judgments, telling whether something is good or bad. Some core thoughts of positivistic and teleological ethics will be formalized in this logic to give a computational model of telling right from wrong. A fictional artificial intelligence is discussed, using insights provided by metamathematics, and the bargaining problem, a philosophical question with wide-ranging implications, will be approached with game theory. But before going into too much detail about how to approach these subjects from a computer scientist’s point of view, it might be helpful to briefly review these ethical topics from a classical point of view.

Three Questions, Six Answers

Kenneth Laudon [6] uses three “critical distinctions” for organizing the literature on ethics, which can be viewed as the three big questions of ethics. The different answers to these questions define a spectrum of ethical thought.

Question One: What is “goodness?”

Phenomenologist’s answer: It is a higher order, and it is given. One must understand the abstract concepts of right and wrong, and act accordingly.

Positivist’s answer: It is whatever we make of it. We have to derive ethical principles for ourselves according to our observations of the real world.

Question Two: Does acting ethically correct mean acting according to certain rules, or acting in a way leading to desired consequences?

Deontologist’s answer: Acting ethically correct, means respecting one’s duties and obligations. Each single act is itself good or bad, regardless of its consequences.

Teleologist’s answer: An action can be judged only by its consequences. Acting right means acting in such a way that the outcome is good. There is no act that is a priori good.

Question Three: What is the scope of morality? Is morality subject to everyone’s individual freedom?

Collectivist’s answer: Ethical standards make sense only if they equally apply to everyone. A rule like “thou shalt not kill,” that was subject to the individual’s approval, would be the equivalent of a rule like “thou shalt not kill, unless it benefits you.”

Individualist’s answer: Nobody should be committed to accepting ethical standards just out of pure principle or we will end up with a morality that people ultimately suffer from.

Given this basic, yet vague understanding of ethical principles, one can develop a more rigorous representation suitable for computation. Such an attempt quickly presents the problem of undecidability, an interesting principle with wide-ranging implications.

Explaining Gödel’s incompleteness theorems in detail is beyond the scope of this article. The interested reader is referred to Kurt Gödel’s original publication [2]. Alternately, Hofstadter’s Pulitzer-Prize-winning classic Gödel, Escher, Bach: An Eternal Golden Braid [5] and Smullyan [8] make the complex ideas behind this topic accessible.

An Infinitely Long Time Ago...

Consider a thought experiment: A computer program, called GURU, shall be designed, sent back through time, and started at a point in time that is infinitely long ago. That program should be capable of telling right from wrong. Such a program would need a set of ground facts, axioms, premises, or whatever term fits one’s philosophy of systematic reasoning. These could be statements telling the machine that something is definitely good or bad. This is not enough, because not every situation the machine will ever be confronted with can be expected to appear in the set of ground facts. Therefore some inference rules or proof techniques would have to be established. The machine would need them to determine whether “compound-situations,” situations described in terms of other situations, are good or bad.

GURU would also need an inference engine that systematically combines simple situations, achieving “compound-situations” based on the inference-rules. Whenever GURU discovers such a situation, it could store it to an infinite memory, called the scroll of all wisdom. If all of the ground facts can safely be assumed, all of the inference rules work, and the inference engine produces correct outputs, then, applying the

ideas of positivistic philosophy, the scroll of all wisdom would indeed have to contain all wisdom after running GURU for an infinitely long time.

implication			
	p	q	$p \rightarrow q$
(1)	true	true	true
(2)	true	false	false
(3)	false	true	true
(4)	false	false	true

act		lie	
	p	q	$p \Rightarrow q$
(1)	bad	bad	bad
(2)	bad	good	good
(3)	good	bad	bad
(4)	good	good	bad

Table 1: A well-known truth-table and a “goodness-table.”

Given such a framework, an actual formalization of the ethical system GURU is supposed to employ can now be considered.

Judging Acts

A set of symbols will be needed, in order to perform systematic reasoning with them. Keep in mind that these symbols are not defined as in mathematical reasoning. These symbols are simply GURU’s machine language, and in that machine language the symbol \rightarrow is not a Boolean negation, and the symbol \overline{p} was chosen only to highlight a polymorphy between Boolean reasoning and ethical judging that will be discussed shortly.

This is the only correct interpretation for GURU’s machine language:

- m is the act of murder.
- \overline{X} are the consequences X has (either good or bad).
- $(\overline{N} \rightarrow \overline{D})$ is an act, where the consequences of not doing it are \overline{N} and the consequences of doing it are \overline{D} .

This is where the teleological approach comes in. GURU judges things in terms of their consequences only. In order to formalize what it means for something to have bad consequences, the symbol bad is defined in terms of the consequences of murder:

\overline{m} are bad.

GURU also needs a way of deciding whether things besides murder are good or bad:

$(\overline{N} \rightarrow \overline{D})$ are good if, and only if, \overline{N} are bad and \overline{D} are good.

That this interpretation is consistent with ethical ideas can be seen, considering the “goodness-table” given in Table 1. The possible situations GURU could be confronted with and the judgement our rule produces will be shown by discussing the table line by line.

1. The consequences of choosing not to do the act are bad. The consequences of choosing to do the act are also bad. This is of course bad. One should try to avoid getting into such a situation.

2. The consequences of choosing not to do the act are bad. The consequences of choosing to do the act are good. This is good. One can choose to do the act, which is good, and one is, in fact, obligated to do it, because the outcomes are bad, otherwise.
3. The consequences of choosing not do the act are good. The consequences of choosing to do the act are bad. There should be no question that this is bad.
4. The consequences of choosing not do the act are good. The consequences of choosing to do the act are also good. It seems a bit counter-intuitive, but this is bad. Let the act of doing S be Sd and the act of not doing S be Sn . By doing Sd one chooses not to do Sn . But the fact that Sn has good consequences, means that one chooses not to do something good. If one knows that something is good, then one is in fact obligated to do it and therefore doing S would violate this principle.

The ability to judge an act by its consequences does not morally justify GURU as being good. It shouldn’t make judgments based solely on rigorous reasoning. It should be primarily committed to doing what’s good, and is therefore also capable of lying, but only if it considers this lie as good.

Judging Judgments

- \overline{Y} are the consequences that GURU says Y has, when it is either telling the truth, or lying.
- \overline{X} are the real consequences X has, and is not subject to any further questioning.
- $(X \Rightarrow \overline{Y})$ is the possible lie that GURU is telling us when asked about X . The consequences of what GURU is telling us are \overline{Y} , while the real consequences of X would be \overline{X} .

Here color was used to indicate the “amount of questioning” necessary for interpreting a symbol.

Table 1 shows the behavior expected from an operation for judging lies.

1. The real consequences of something are bad. GURU is lying, telling us something that has bad consequences. Such a lie is bad. The fact that the real consequences would have been bad as well doesn’t change anything about the fact that the lie has bad consequences.
2. The real consequences of something are bad. GURU is lying, telling us something that has good consequences. Such a lie is good. The real consequences would have been bad, but by lying, GURU has turned this situation into one that has good consequences.
3. The real consequences of something are good. GURU is lying, telling us something that has bad consequences. There should be no doubt that this is bad.
4. The real consequences of something are good. GURU is lying, telling us something that has good consequences. This is bad, because GURU shouldn’t lie when there is no need to do so.

The consequences of a lie can be defined as follows:

$(\overline{X \Rightarrow \overline{Y}})$ is good if, and only if is $(\overline{X \Rightarrow \overline{X}}) \rightarrow \overline{Y}$ on the scroll of all wisdom.

The equivalence between a lie $X \Rightarrow \overline{Y}$ and the act of lying $\overline{X} \rightarrow \overline{Y}$ is demonstrated in Table 1.

The problem is determining the true consequences of $X(\overline{X})$, if all GURU ever does is either tell the truth or lie about $X(\overline{X})$. Here GURU has to question itself. If it suspects \overline{X} to be a lie, and it knows that it only makes lies that have good consequences, it checks whether the lie $X \Rightarrow \overline{X}$ would be good. Substituting this for X gives the definition $(X \Rightarrow \overline{Y}) \equiv (X \Rightarrow \overline{X}) \rightarrow \overline{Y}$.

Such a definition wouldn't lead us anywhere, because it would be circular, but, since GURU's memory contains all wisdom, $(X \Rightarrow \overline{Y})$ can simply be defined in terms of $(X \Rightarrow \overline{X}) \rightarrow \overline{Y}$ being on the scroll of all wisdom, which is not per se circular.

Based on these rules, GURU is now capable of telling good from bad, and even telling whether judgments about what is good or bad are themselves good or bad. But does this say anything about whether GURU is good?

When Gödel Meets GURU

In fact, the rules from the previous sections were taken from Smullyan [8] and are a subset of the formal system he presents as *Craig's machine*. The reader interested in its formal details is referred to his work for proof that this system is in fact Gödelian. I built the GURU story around Craig's machine only to provide an interpretation for it, and to show how reasoning could work on an ethical level, but a machine or a metamathematician is not interested in anything else but these four rules in their purely symbolic form.

The Gödel sentence G such that $G \equiv \neg PG$ is a correct sentence of a formal system that states "you cannot prove me within my system." Note that some of the details were left out here. PG could in terms of GURU's "machine-language" be an operation finding out whether G is on the scroll of all wisdom. One would further have to define the concepts of \equiv and \neg in terms of GURU's machine-language and would have to know some more details about how GURU really operates. For Craig's machine, Smullyan shows that $(m \Rightarrow \overline{m}) \rightarrow \overline{m}$ is such a sentence.

In terms of GURU, the Gödelian sentence $(m \Rightarrow \overline{m}) \rightarrow \overline{m}$ would have an interpretation like, "What are the consequences of an act, where the consequences of doing this act are the same as the consequences of murder, and the consequences of not doing this act are that I suspect you of lying when you say that murder is bad?"

Although this question is a skeptical way of asking about the most basic fact, "Is murder good or bad?", it is not possible for GURU to answer it.

Divide The Dollar!

In the previous sections, the first two ethical questions were considered. They were handled together because they are similar, but the third question, in contrast, is widely unrelated to the other ones. A fresh approach will be taken for this question.

The **divide-the-dollar** situation is one of the philosophical classics. Two people are given a dollar, but only under the condition that they find an agreement on how to divide it. This problem is representative of the collectivist/individualist-debate, yet very simple. It is in fact so simple that it is frequently found in elementary school math-

textbooks, in wordings like "This year Grandma gave Mary and John one dollar for Christmas. How much money does each of them get?" Part of the naive approach an elementary school pupil takes when confronted with this problem is the assumption that each of the two bargainers ought to get the same amount of money.

Such a standard can only work if all the individuals involved in the bargaining-process accept it without further questioning. Someone who doesn't accept it could demand 99 cents, arguing the following way: "If I reject every possible bargain, except for the one that leaves me with 99 cents, then my opponent has only two options left: either accepting or rejecting this one-and-only offer. Accepting the offer leaves the opponent with one cent, otherwise neither of us will get anything. This is why the opponent has to accept."

From a collectivistic standpoint, this seems completely nonsensical, because how can one individual justify this in such a way that this ethical justification wouldn't apply if the other individual made the same argument? (This principle frequently appears in the ethical literature as reciprocity.) Note that the elementary-school approach presented above is already a collectivistic one, because it is subject to a global ethical standard. From an individualistic standpoint, on the other hand, the collectivist's commitment to the principle of reciprocity itself seems nonsensical. What motivation drives a bargainer to accept the 50-cents-deal, if 99 cents would also be possible. Given this concept, one can now go on to generalize the idea, so it applies not only to thought-experiments.

About Idealized Rational Individuals

From a strong individualistic view, consider the following statement:

Given any specific situation an individual might find himself in, it will benefit the most, in the long run, by making such a decision that maximizes its own personal benefit.

An example where the above statement does not hold was already given. Two individualists will not come to any agreement in the divide-the-dollar situation, which is why they won't get any money. As far as ethics is concerned, we are done at this point, and are forced to draw back to a weaker version of the individualistic view, like "there are situations, where an individual...."

The bargaining-game, as described by John F. Nash [7], provides some more insight, not by proving any ethical views, but by saying something about how common these divide-the-dollar situations really are. It doesn't necessarily take two "hardcore-individualists" confronted with a situation as theoretical and as made-up as the divide-the-dollar problem, to find a counter-example for the individualistic view. It completely suffices to let two idealized rational individuals take part in a two-person game, in which they are supposed to find an agreement on trading some of the goods they possess, without the use of money (which would only be a special case of the situation considered).

Nash's Approach To The Bargaining Problem

Nash uses the example shown in Table 2 to illustrate a bargaining situation. The table shows goods, some of which belong to Bill and some of which belong to Jack, and their utilities to Bill and Jack. (Nash uses the notion of utility to formalize the concept behind the value of a good. The fact that goods can be of differing value for different people could be seen as the major driving force behind economy.) Bill and

Bill's goods	Utility to Bill	Utility to Jack
book	2	4
whip	2	2
ball	2	1
bat	2	2
box	4	1
Jack's goods	Utility to Bill	Utility to Jack
pen	10	1
toy	4	1
knife	6	2
hat	2	2

Table 2: Nash's bargaining example [7].

Jack are supposed to trade them without the use of money, or any other common exchange medium.

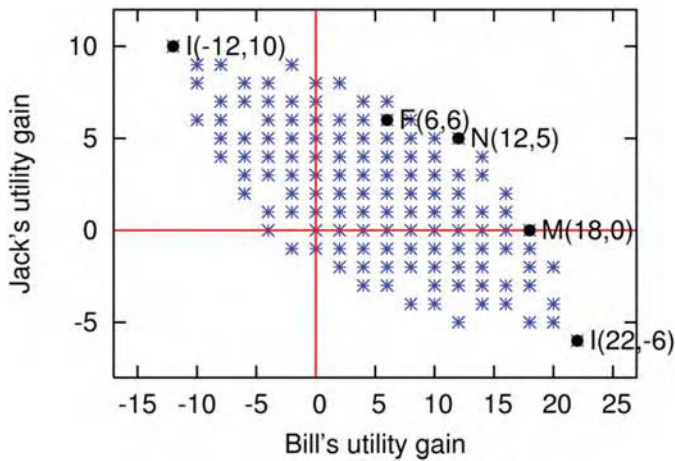


Figure 1: Visualization of Nash's original bargaining example [7].




Figure 2: Visualization of another case.

Nash visualized the situation in a plot as shown in Figure 1. The idea is that every possible solution to the problem will plot to a point (x, y) in the plane. Here x is the gain in utility the first individual could expect from that solution and y is the gain in utility the second individual could expect.

One possible solution to Nash's example could be that Bill gives Jack the book and nothing else. It plots to $(-2, 4)$, denoting that Bill loses two utility units and Jack gains four. This solution can directly be looked up in Table 2, but there is no need to restrict Bill and Jack to trading only one of the goods. Bill and Jack could also agree to a trade like Bill giving Jack the book and Jack giving Bill the knife. Here Bill loses two utility units, because he has to give away the book, and he gains six, because he gets the knife, leaving him with an overall gain of four. Jack loses two, because of the knife, and gains four, because of the book, leaving him with an overall gain of two. This solution would then plot to $(4, 2)$. Another solution could be that Jack gives everything to Bill, which plots to $I(22, -6)$.

$I_{\text{Bill}}(b, j)$	b is a minimum
$I_{\text{Jack}}(b, j)$	j is a minimum
$F(b, j)$	$b + j$ is the minimum also satisfying $b = j$
$M(b, j)$	$b + j$ is a maximum
$N(b, j)$	$b * j$ is a maximum

Table 3: Possible solutions to a bargaining problem.

If Bill and Jack's primary motivation were to find a "fair" trade, they could go for a solution in $F(6, 6)$. For example, Bill could give Jack the book, whip, bat, and box, and Jack could give Bill the pen and the knife. In this case, they would both gain the same amount of utility, which is six. The two of them would then be 12 utility units better off than before. (Here, a global measure of utility is used, obtained by arithmetically adding the two individual's utilities.) But can they do better?

If they wanted to gain as much "global utility" as possible, they could go for a solution in $M(18, 0)$, which would even have a global value of 18 utility units. One of these would be that Bill gives Jack the book, and Jack gives Bill the pen, toy, and knife. Probably Jack would not like the fact that he has to give away all of these items, and not gain anything for himself. Still, from a collectivist standpoint, it is clearly a better solution, because if Bill has more use for the items, then why should Jack possess them?

The solution suggested by Nash [7] is very interesting. It can be found in $N(12, 5)$. In this case, Bill gives Jack the book, whip, ball, and the bat, and Jack gives Bill the pen, toy, and the knife. This is the outcome Nash expects, when two idealized rational individuals bargain. Here it is assumed that the individual who has more potential to benefit will also be the stronger one in the bargaining process.

Figure 2 shows another possible bargaining situation and Table 3 summarizes the possible solutions, and their criteria. These criteria are the results of different philosophies applied to the same problem, and may help to show some of the concrete impact, that concepts as abstract as individualistic and collectivistic ethics have on everyday problems like trading.

Conclusion

Question three was approached by considering bargaining problems. That the maximum "net wealth" of a group of people is not necessarily reached, when each individual tries to maximize its own wealth was shown by considering the divide-the-dollar problem. The discussion was then extended to show that similar principles apply for the more general formulation of bargaining situations as studied by Nash.

Question two was handled here indirectly in the interpretation of GURU's principles of reasoning, and gets most obvious, possibly when examining the situations where GURU is supposed to lie. One might want to program GURU in such a way that, out of principle, it doesn't lie. As a result, GURU wouldn't be a Gödelian machine any more, but this would clearly be on the deontologists' side of the philosophical spectrum.

Question one is answered by Gödel himself quite clearly. In his posthumously published *Philosophical Essays* [3] he states:

I am under the impression that [...] the Platonistic view is the only one tenable. Thereby I mean the view that mathematics describes a non-sensual reality, which exists independently both of the acts and the dispositions of the human mind and is only perceived, and probably perceived very incompletely, by the human mind [3].

References

1. Encyclopedia Britannica ME. 1999. CD-ROM. *Normative Ethics*.
2. Gödel, K. 1931. Über formal unentscheidbare sätze der principia mathematica und Verwandter Systeme I. *Collected Works*, vol. 3. S. Feferman, Ed. Oxford University Press.
3. Gödel, K. 1995. Some basic theorems on the foundations of mathematics and their philosophical implications. *Unpublished Philosophical Essays*, F. A. Rodríguez-Consuegra, Ed., chap. 2. Birkhäuser Verlag. 144-147.
4. Guerrierio, G. 2002. *Kurt Gödel, Logische Paradoxien und Mathematische Wahrheit*. Spektrum der Wissenschaft Verlagsgesellschaft.
5. Hofstadter, D. R. 1999. *Gödel Escher Bach: An Eternal Golden Braid*. Basic Books, Inc.
6. Laudon, K. C. 1995. Ethical concepts and information technology. *Comm. of the ACM* 38, 12. 33-39.
7. Nash, J. F. 1950. The bargaining problem. *Econometrica* 18.
8. Smullyan, R. 1987. *Forever Undecided*. Oxford University Press.

Biography

Richard Bergmair (rbergmair@acm.org) is a final-year undergraduate student of Computer Science at the University of Derby in Austria. He completed a five-year program in Computer Science at the level of secondary education and has been doing Software Engineering for IBM since he was 14. His primary research interests are Computational Linguistics and logic programming.

This article originally appeared in *Crossroads* 10.3 (Spring 2004), "Ethics and Computer Science."

Visit the NEW Crossroads site at www.acm.org/crossroads

INNOVATION



BIG-COMPANY REACH, SMALL-COMPANY ATTITUDE.

AAI Corporation's innovative technologies include unmanned aircraft, training systems, automated test and maintenance equipment, armament systems, aviation ground support equipment, and cutting-edge service solutions.

Headquartered in Hunt Valley, Md., AAI offers rewarding career opportunities for engineers and business professionals in all disciplines, including:

- Varied, employee-driven advancement tracks
- Full tuition reimbursement for continuing education
- In-house professional development courses

Voted one of Baltimore Business Journal's "Best Places to Work" in 2006 & 2007, AAI provides an exceptional benefits package for full-time employees. Visit www.aaicorp.com to apply. EOEM/F/D/V



AAI
TEXTRON Systems
INNOVATION THAT WORKS.
aaicorp.com

Innovation That Works is a registered trademark of AAI Corporation.